

Class07: Clustering and PCA


Eric Garcia Ordunez (A16316409)

#Clustering

First let's make some data cluster so we can get a feel for these methods and how to work with them.

We can use the `rnorm` function to get random numbers from a normal distribution around a given `mean`.

```
hist(rnorm(5000, mean= 3))
```



Let's get 30 points with a mean of 3 and another 30 with a mean of -3.

```
tmp<- c(rnorm(30, mean=3), rnorm(30, mean = -3))
tmp
```

```
[1] 3.9432850 4.8722022 2.9749674 4.5421338 3.7049116 3.4127286
[7] 3.2186652 3.4632978 2.7982153 2.9033447 0.6614982 2.7343330
[13] 3.6429153 5.0387579 3.0197487 3.7484452 2.2531125 2.7423207
[19] 1.6786133 2.8700320 3.5433802 3.0470769 3.9144089 6.3532483
[25] 3.8134372 0.5727041 3.1380112 2.5821383 2.5918790 1.8088294
[31] -2.3552644 -3.5743954 -4.0743972 -2.3015844 -2.1938405 -2.0172069
[37] -2.4222090 -2.4466368 -2.5584100 -1.7113266 -3.0537078 -2.5957530
[43] -1.1300887 -1.7125178 -2.6453741 -4.5353387 -3.0582564 -2.7490909
[49] -2.7472791 -2.9400898 -1.6865771 -2.6601737 -5.0071821 -2.5816637
[55] -2.7327963 -0.7433062 -3.4543034 -1.7691234 -1.8400599 -5.5580137
```

Put two these together:

```
x<- cbind(x= tmp, y=rev(tmp))
x
```


```
      x      y
[1,] 3.9432850 -5.5580137
[2,] 4.8722022 -1.8400599
[3,] 2.9749674 -1.7691234
[4,] 4.5421338 -3.4543034
[5,] 3.7049116 -0.7433062
[6,] 3.4127286 -2.7327963
[7,] 3.2186652 -2.5816637
[8,] 3.4632978 -5.0071821
[9,] 2.7982153 -2.6601737
[10,] 2.9033447 -1.6865771
[11,] 0.6614982 -2.9400898
[12,] 2.7343330 -2.7472791
[13,] 3.6429153 -2.7490909
[14,] 5.0387579 -3.0582564
[15,] 3.0197487 -4.5353387
[16,] 3.7484452 -2.6453741
[17,] 2.2531125 -1.7125178
[18,] 2.7423207 -1.1300887
[19,] 1.6786133 -2.5957530
[20,] 2.8700320 -3.0537078
[21,] 3.5433802 -1.7113266
```

```

[22,] 3.0470769 -2.5584100
[23,] 3.9144089 -2.4466368
[24,] 6.3532483 -2.4222090
[25,] 3.8134372 -2.0172069
[26,] 0.5727041 -2.1938405
[27,] 3.1380112 -2.3015844
[28,] 2.5821383 -4.0743972
[29,] 2.5918790 -3.5743954
[30,] 1.8088294 -2.3552644
[31,] -2.3552644 1.8088294
[32,] -3.5743954 2.5918790
[33,] -4.0743972 2.5821383
[34,] -2.3015844 3.1380112
[35,] -2.1938405 0.5727041
[36,] -2.0172069 3.8134372
[37,] -2.4222090 6.3532483
[38,] -2.4466368 3.9144089
[39,] -2.5584100 3.0470769
[40,] -1.7113266 3.5433802
[41,] -3.0537078 2.8700320
[42,] -2.5957530 1.6786133
[43,] -1.1300887 2.7423207
[44,] -1.7125178 2.2531125
[45,] -2.6453741 3.7484452
[46,] -4.5353387 3.0197487
[47,] -3.0582564 5.0387579
[48,] -2.7490909 3.6429153
[49,] -2.7472791 2.7343330
[50,] -2.9400898 0.6614982
[51,] -1.6865771 2.9033447
[52,] -2.6601737 2.7982153
[53,] -5.0071821 3.4632978
[54,] -2.5816637 3.2186652
[55,] -2.7327963 3.4127286
[56,] -0.7433062 3.7049116
[57,] -3.4543034 4.5421338
[58,] -1.7691234 2.9749674
[59,] -1.8400599 4.8722022
[60,] -5.5580137 3.9432850

```

```
plot(x)
```



##K-means clustering.

Very popular clustering method that we can use with the `kmeans()` function in base R.

```
km<- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.695199	3.186288
2	3.186288	-2.695199

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 72.94783 72.94783
(between_SS / total_SS = 87.7 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q How many points are in each cluster?

```
km$size
```

```
[1] 30 30
```

Q what component of your result object details - cluster size?

- cluster assignment/membership?

```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


- cluster center?

```
km$centers
```

```
      x      y
1 -2.695199  3.186288
2  3.186288 -2.695199
```


Q. Plot x colored by the kmeans cluster assignment and add cluster centers blue

```
mycols <- c(1,2)
plot(x, col=km$cluster)
points(km$centers, col= "blue", pch=15, cex= 3)
```



Q. Let's cluster into 3 groups or same x data and make a plot.

```
km <- kmeans(x, centers=3)
plot(x, col= km$cluster)
```



Hierarchical Clustering

We can use the `hclust()` function for Hierarchical Clustering.

Unlike `kmeans()`, where we could just pass in our data as input, we need to give `hclust` a “distance matrix”.

We will use the `dist()` function to start with.

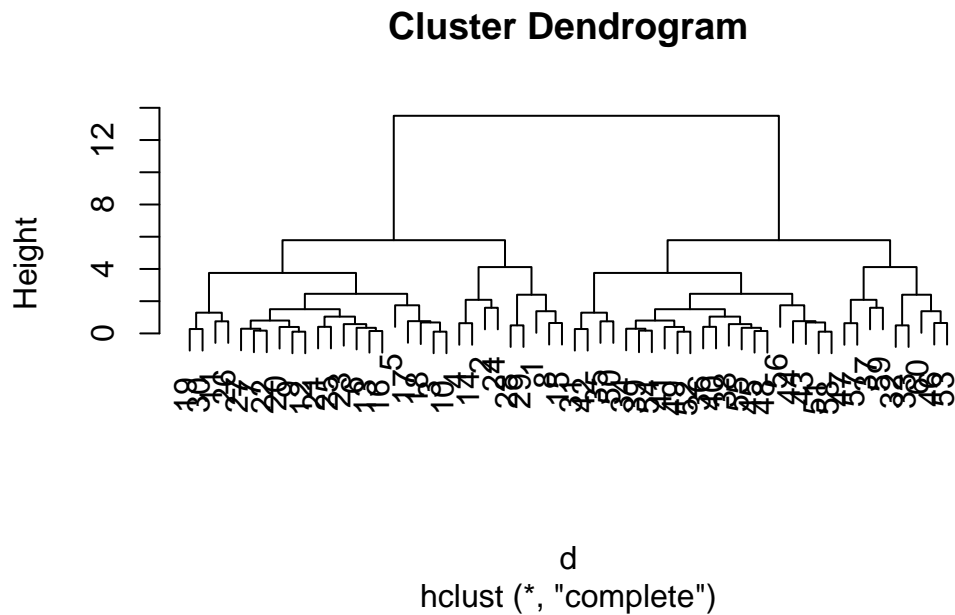
```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

```
plot(hc)
```



I can now “cut” my tree with the `cutree()` to yield a cluster membership vector.

```
grps <- cutree(hc, h=8)
grps
```


```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

You can also tell `cutree()` to cut where it yields “k” groups.

```
cutree(hc, k=2)
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(x, col= grps)
```

Principle Component Analysis (PCA)

```
# Adding the `row.names()` helps the first column to have its own.
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names= 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674

Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41


Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  4
```

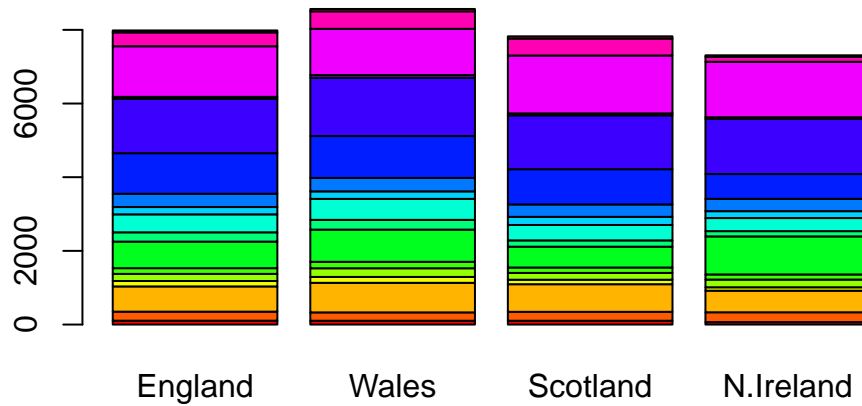
Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



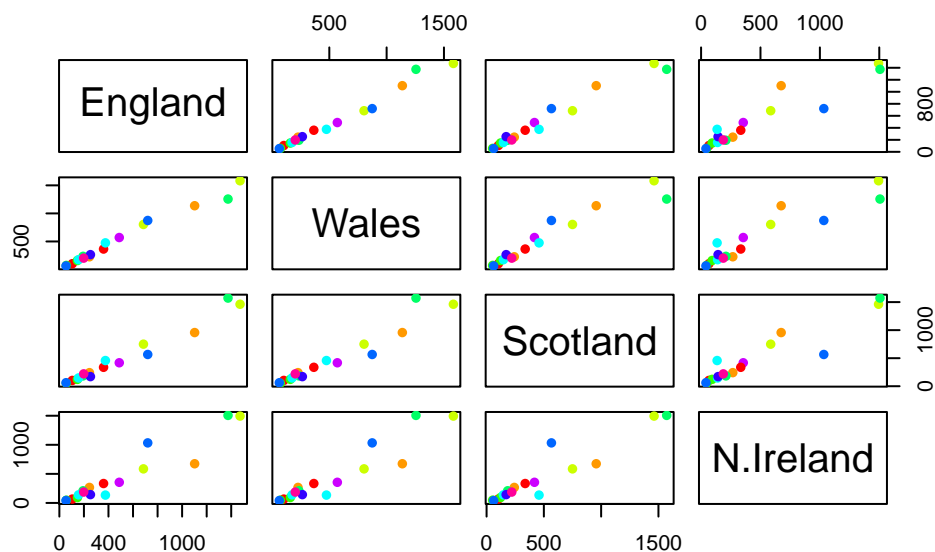
Q3: Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```



Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The main PCA function in base R is called `prcomp()` it expects the transpose of our data.

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

```
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "darkgreen"), pch=16)
```

